

RecSys 2017, Como, Italy

entity2rec

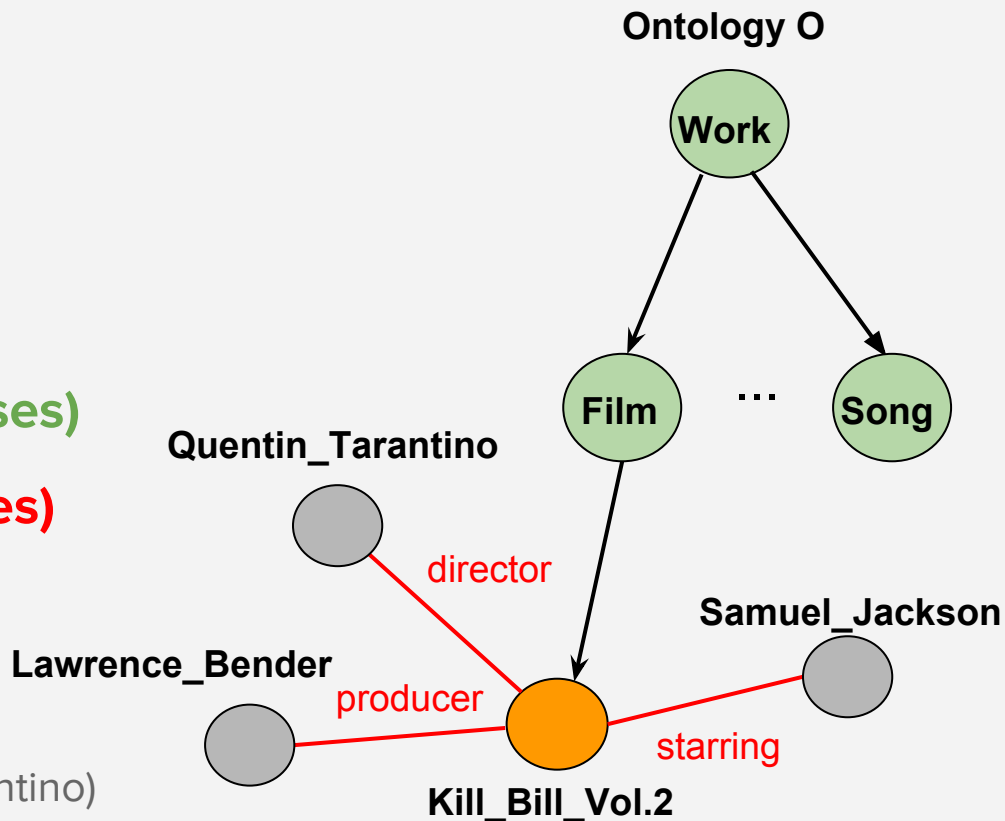


Learning User-Item Relatedness
from Knowledge Graphs for Top-N
Item Recommendation

Enrico Palumbo, Giuseppe Rizzo, Raphaël Troncy

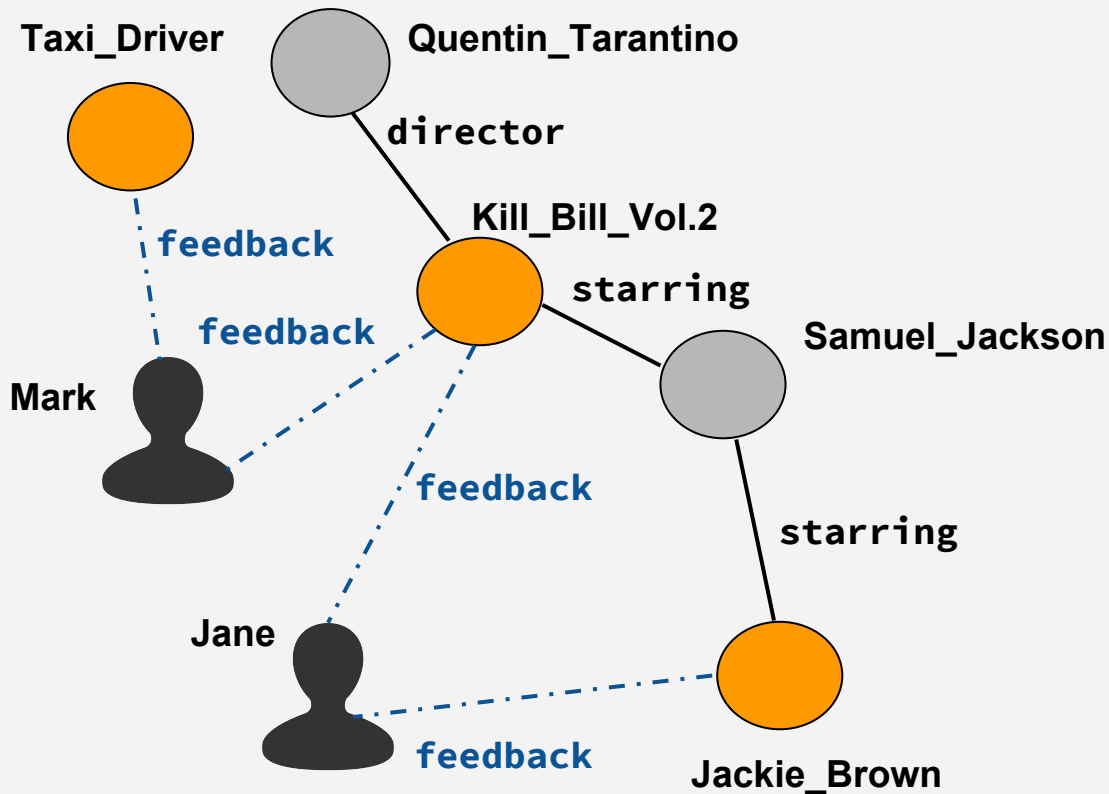
Knowledge Graph

- $K = (E, R, O)$
- E = entities
- O = ontology
- O defines **entity types Λ (classes)** and **relation types Γ (properties)**
- $O: \Lambda \rightarrow \Gamma$
- $R \subseteq E \times \Gamma \times E$
- (Kill_Bill_Vol.2, starring, Quentin_Tarantino)



Knowledge Graph for RS

- $U \subset E$ = users
- $I \subset E \setminus U$ = items
- 'feedback' property:
(u, feedback, i)
- Collaborative and
content information
- **Hybrid** recommender
systems

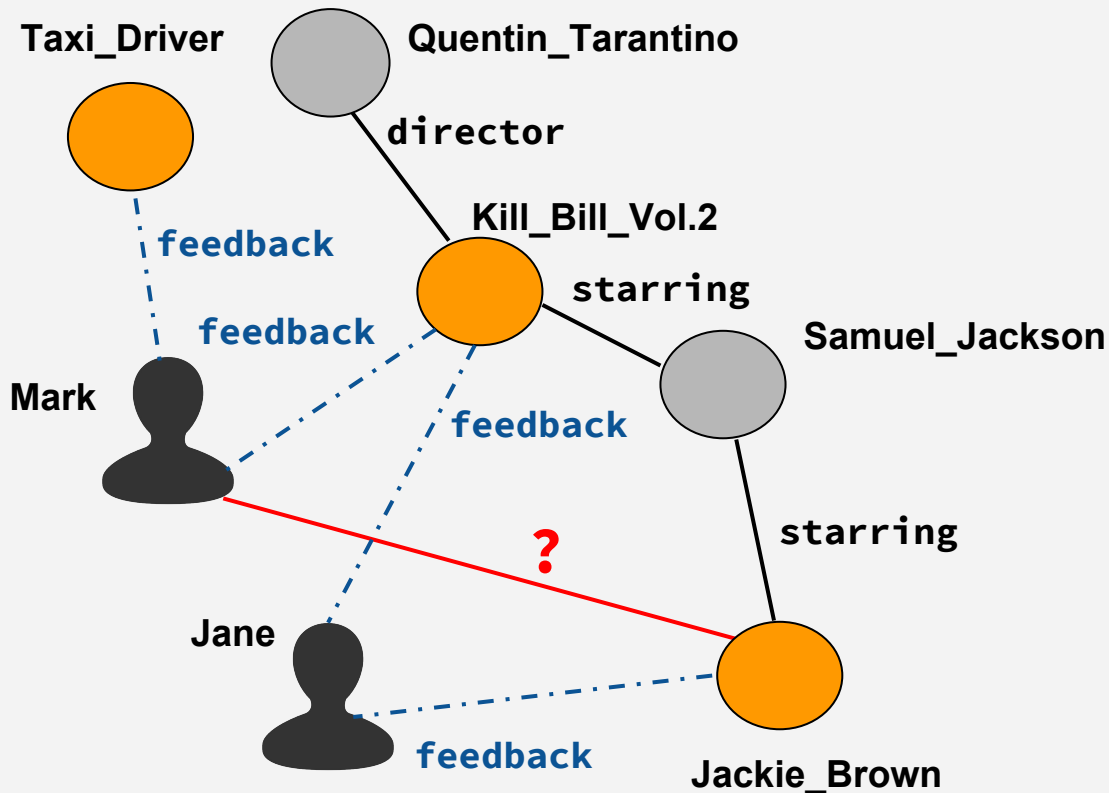


Objective

Learn **user-item relatedness** from the **Knowledge Graph** for top-N **item recommendation**.


Learn entity features encoding:

- **Graph structure:** tightly connected entities should be more related
- **Semantics:** not all properties have the same importance

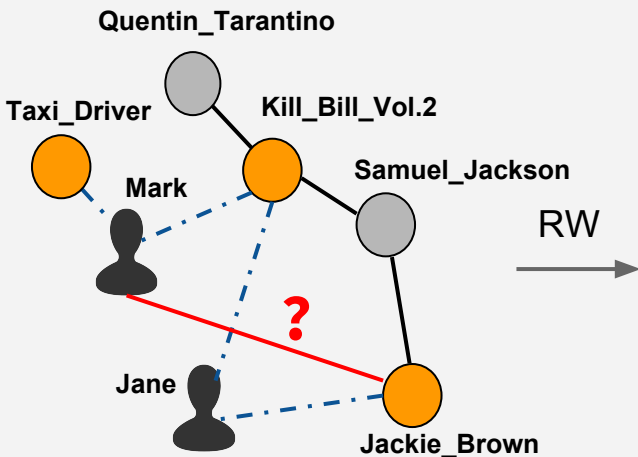


Graph structure: node2vec

- **Feature learning** from networks
- Adaptation of **word2vec** on **graph structures** using **random walks**
- Maps **nodes** in a graph into an **euclidean space** preserving the **graph structure**



Graph structure: node2vec



Graph

RW

Mark, Kill_Bill_Vol.2,
Quentin_Tarantino, ...

Jane, Kill_Bill_Vol.2,
Jurassic_Park ...

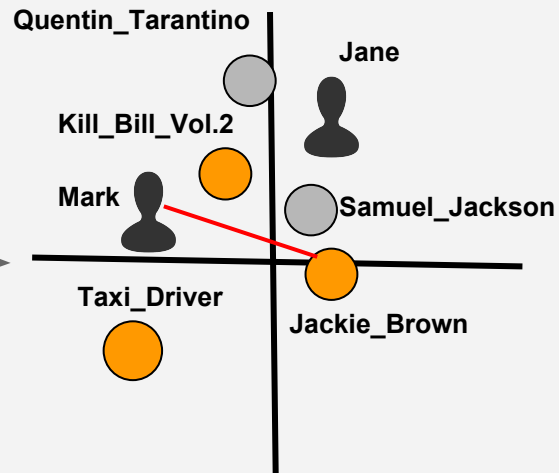
Kill_Bill_Vol.2, Samuel_Jackson,
Jurassic_Park, Jane...

Samuel_Jackson, Jurassic_Park,
Kill_Bill_Vol.2, ...

...

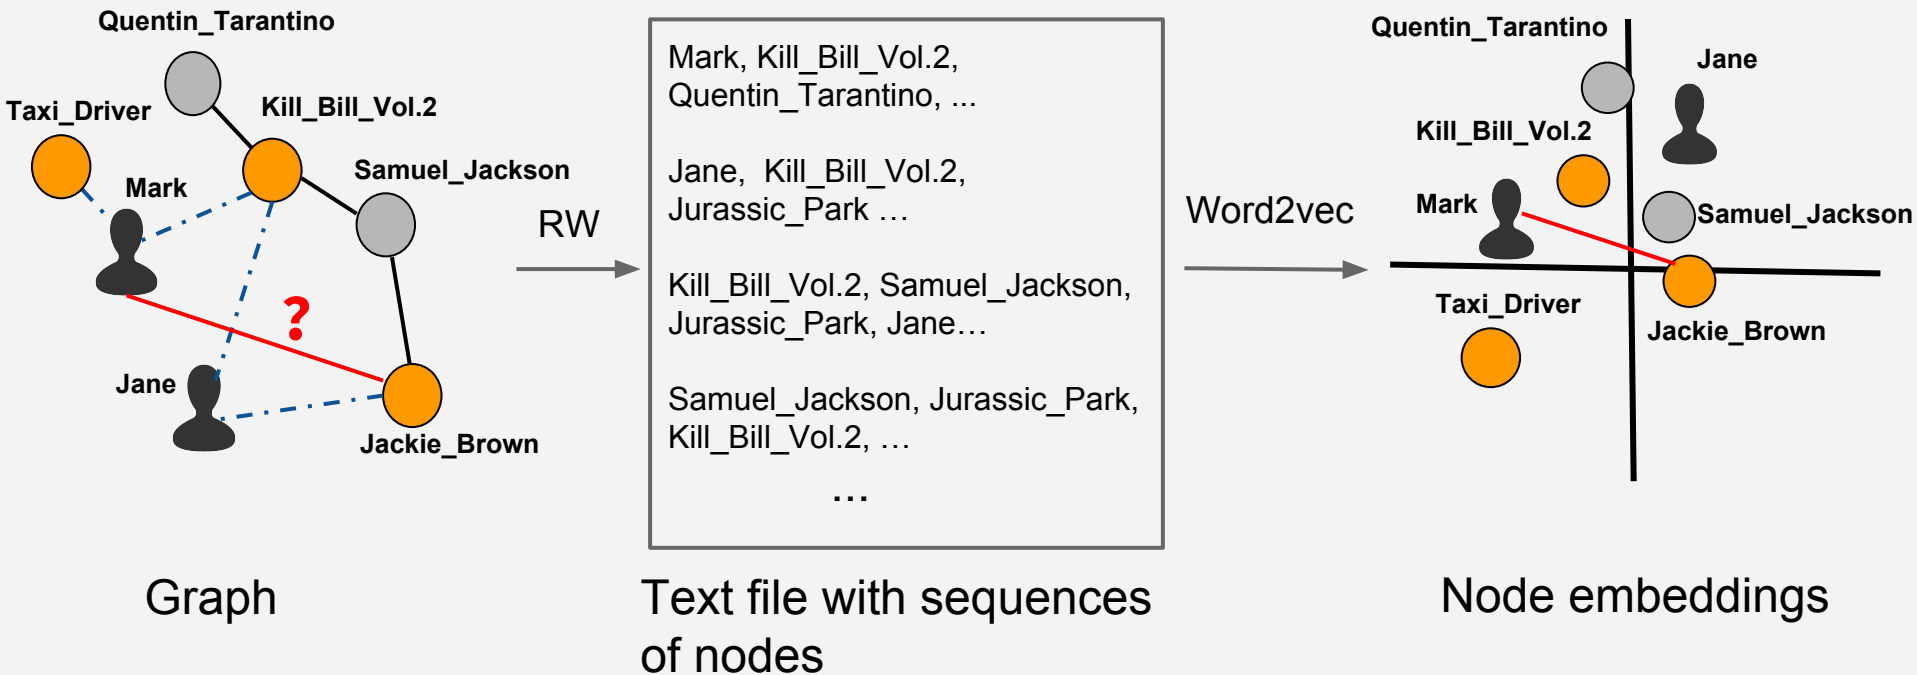
Text file with sequences
of nodes

Word2vec




Node embeddings

Graph structure: node2vec




No semantics, not optimized for recommendations

Semantics: property-specific relatedness



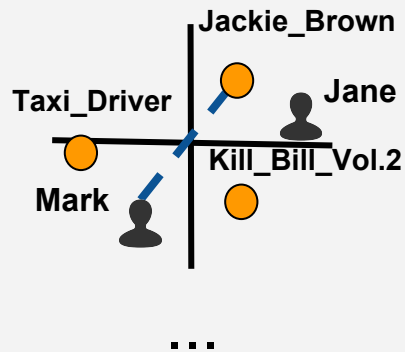
Collaborative filtering: $\rho_{\text{feedback}}(u, i) = d(x_{\text{feedback}}(u), x_{\text{feedback}}(i))$

Semantics: property-specific relatedness



Content filtering: $\rho_{\text{starring}}(u,i) = D (x_{\text{starring}}(i'), x_{\text{starring}}(i))$

Learning to rank



$\rho_{\text{feedback}}(u,i)$

$\rho_p(u,i)$

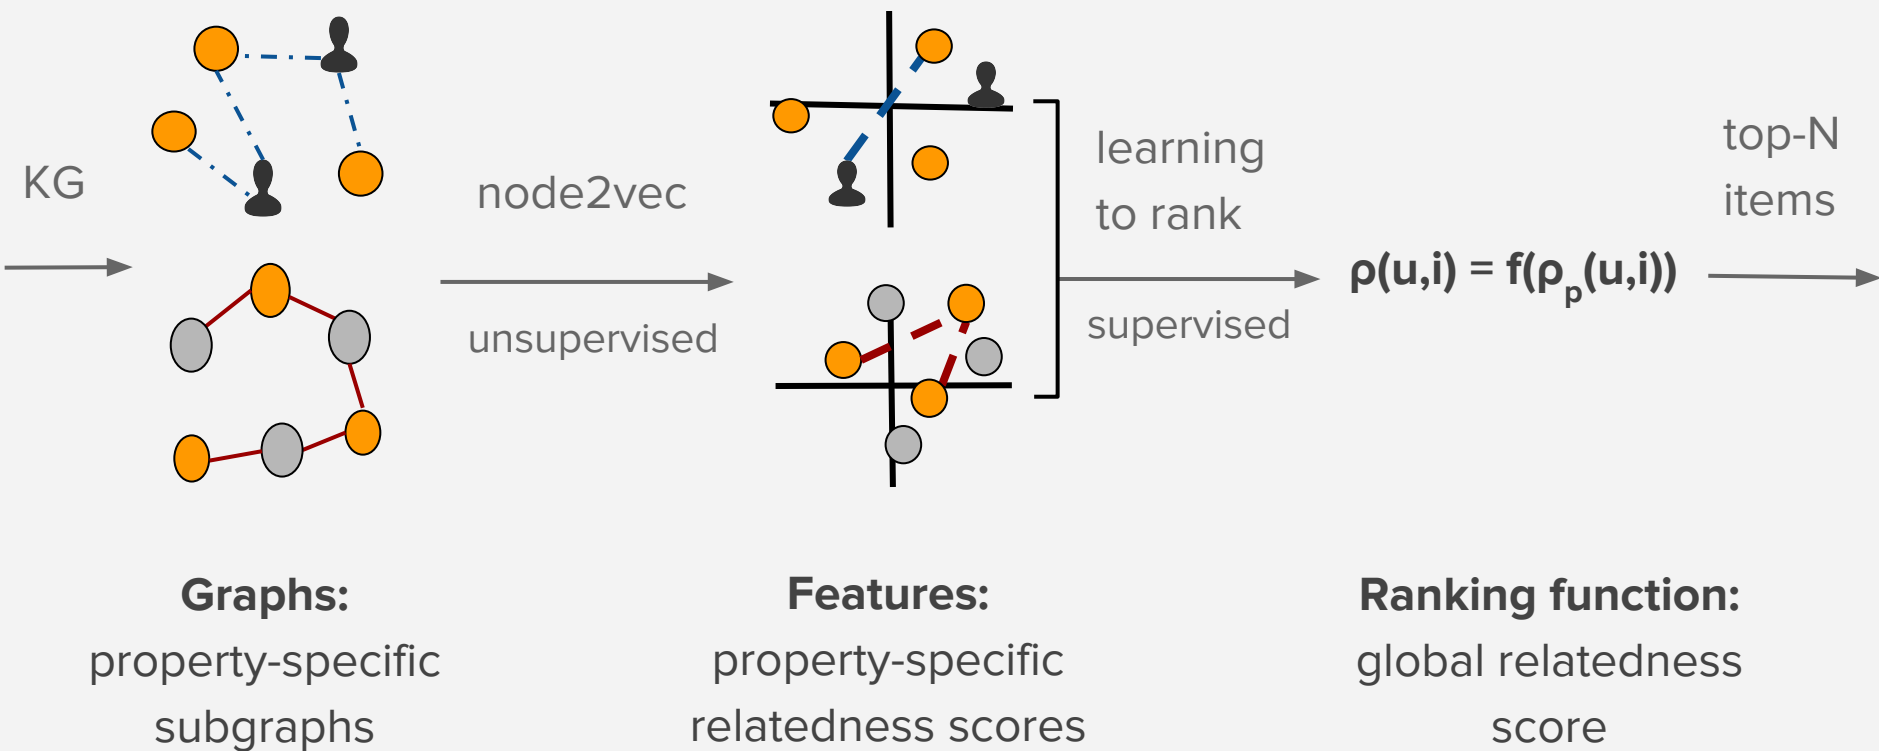
$\rho_{\text{starring}}(u,i)$

features

Learning to rank: $\rho(u,i) = f(\rho_p(u,i))$

- Global user-item relatedness ρ as a **ranking function**
- **List-wise approach**: optimizing directly ranking with information retrieval metrics (e.g. P@N)

Approach



Experimental setup

- **Movielens_1M:** 1,000,209 ratings, 3900 movies, 6040 users
- **DBpedia:** we leverage DBpedia mappings of Movielens items to build the knowledge graph. We use the DBpedia Ontology to define properties of the class 'Film'. Some items are missing: **948978 ratings, 3225 movies, 6040 users.**
- **Evaluation protocol:** train, val test set split. Add 100 'false candidates' to each user in the test set to simulate 'real' ranking scenario.
- **Baselines:** popular collaborative filtering algorithms: NMF, SVD, ItemKNN, MostPopular

Results

Model	P@5	P@10	MAP
entity2rec	0.2814	0.2127	0.4232
Most Popular	0.2154	0.1815	0.2907
NMF	0.1208	0.1150	0.1758
SVD	0.0543	0.0469	0.0888
ItemKNN	0.0463	0.0232	0.0990

entity2rec and baselines evaluated on the test set

Feature selection

Features	P@5	P@10	MAP
ρ_{feedback}	0.2317	0.1708	0.3550
$\rho_{\text{dbo:director}}$	0.0219	0.0211	0.0949
$\rho_{\text{dbo:producer}}$	0.0119	0.0241	0.0498
$\rho_{\text{dbo:starring}}$	0.0128	0.0372	0.0728
$\rho_{\text{dct:subject}}$	0.0285	0.0326	0.0688
$\rho_{\text{dbo:writer}}$	0.0061	0.0216	0.0472
entity2rec	0.2814	0.2127	0.4232

Performance using one property at the time in the learning to rank model vs global model

Conclusions

- Novel approach to learn user-item relatedness that leverages the **semantics** and the **graph structure** of the **knowledge graph** to provide **recommendations**
- **Hybrid** approach, includes both collaborative and content information
- **No feature engineering**, using an unsupervised feature learning stage to learn entity embeddings and derive property-specific relatedness scores and a supervised learning to rank stage to combine them
- **High ranking precision** on the Movielens dataset
- **No black box effect**, features have a clear interpretation and can be **easily configured** to a particular recommendation need
- **Parallelizable** and easily adaptable to an **online learning** context

Future work

- **Test** against **more competitive baselines** and on **more domains** (e.g. music, books, tourism...)
- New experiments to answer to questions such as:
 - what is the importance of content based features and how does it vary with different degrees of data sparsity?
 - is the unsupervised learning stage or the supervised learning to rank stage that affects the performance the most?
 - what is the best approach to adapt the system to an online learning context?
- **Parallelized** and adapted to **an online learning context**
- **Software implementation** has to be further **tested** and **extended**¹

1: <https://github.com/MultimediaSemantics/entity2rec>

Thank you!

- Enrico Palumbo, Data Scientist, ISMB, PhD Student, Eurecom - PoliTo
- Mail: palumbo@ismb.it, enrico.palumbo@eurecom.fr
- Personal page: www.enricopal.github.io
- Slides: <http://www.slideshare.net/EnricoPalumbo2>
- Github: www.github.com/enricopal, www.github.com/MultimediaSemantics
- Twitter: <https://twitter.com/enricopalumbo91>



Turin, Italy, www.ismb.it



Sophia Antipolis, France, www.eurecom.fr

The logo for the PasTime project features the word 'PASTIME' in a large, bold, sans-serif font. The letter 'T' is stylized with a yellow outline and a black fill.

PasTime project


Graph structure: node2vec

Learning vector representation of nodes preserving the **local neighborhood** in the feature space.

Flexible definition of node **neighborhood** based on 2nd order **random walk** that is sensitive to different **connectivity** patterns.

p: return parameter

q: in-out parameter



Transition probability

Hyper-parameters optimization

P,Q	Learning to Rank	P@5	P@10	MAP
1,4	LambdaMart	0.0791	0.0293	0.1717
4,1	LambdaMart	0.0182	0.0193	0.0570
1,1	LambdaMart	0.0174	0.0188	0.0565
1,4	AdaRank	0.0134	0.0098	0.0278
4,1	AdaRank	0.0078	0.0083	0.0286
1,1	AdaRank	0.0109	0.0098	0.0358

Grid-search of learning to rank algorithm and (p,q) node2vec hyper-parameters on a validation set

Learning to Rank

- **AdaRank:** Xu, Jun, and Hang Li. "Adarank: a boosting algorithm for information retrieval." Proceedings of the 30th annual international ACM SIGIR conference on Research and development in information retrieval. ACM, 2007.
 - Boosting algorithm, AdaBoost -> Ranking
 - Start from weak rankers and iteratively combines them to correct errors of the previous weak ranker
 - As a weak ranker they use the feature that has the optimal weighted ranking performance
- **LambdaMart:** Burges, Christopher JC. "From ranknet to lambdarank to lambdamart: An overview." Learning 11.23-581 (2010): 81
 - LambdaRank: directly defines the gradients without using the cost function
 - MART: boosting regression trees
 - LambdaMART: LambdaRank using MART as a ranking function

Baselines

- **ItemKNN with baselines:** Yehuda Koren. *Factor in the neighbors: scalable and accurate collaborative filtering*. 2010. URL: <http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf>
- **SVD:** Yehuda Koren, Robert Bell, and Chris Volinsky. *Matrix factorization techniques for recommender systems*. 2009.
- **NMF:** Xin Luo, Mengchu Zhou, Yunni Xia, and Qinsheng Zhu. *An efficient non-negative matrix factorization-based approach to collaborative filtering for recommender systems*. 2014.
- **Python library:** <http://surpriselib.com/>